

# ECE 20875 Final Project - Path 3


**Team Members:** Heng-I Chu, Xinyu Liu

**Purdue Usernames:** chu244, liu3680


**Github Repository:** <https://github.com/lvorchu/Data-Security-in-Model-Training.git>

## Dataset

In this project, we will be working with the scikit-learn digits dataset, which consists of 8x8 grayscale images of handwritten digits 0 through 9. The images were processed by flattening each 8x8 data into a vector with a size of 64, where each element represents the intensity of a pixel from 0 to 16. The dataset contains 1797 samples in total, that is evenly distributed across digits 0 through 9.



Print out and plot the numbers of the class [2, 0, 8, 7, 5]:



## Analysis and Model Choices

We will train three different models: Gaussian Naive Bayes, K-Nearest Neighbors, and Multilayer Perceptron, using 40% of the total dataset. We will then test the accuracy of the three models using 60% of the total dataset and 100% of the dataset.

### GaussianNB

Gaussian Naive Bayes is a simple yet powerful probabilistic classifier based on Bayes' Theorem. The model assumes that the feature values follow a Gaussian distribution for each class, and the probability of the labels can be predicted by combining multiple Gaussian distributions of the feature values.

### KNN

K-Nearest Neighbors is a non-parametric learning algorithm that classifies a data point by finding the majority label among its k closest training samples in feature space.

### MLP

Multilayer Perceptron is a model that simulates a neural network, consisting of an input layer, one or more hidden layers, and an output layer. Each layer contains neurons that contain linear and nonlinear components, allowing the model to learn complex relationships in the data.

## Results on Normal Dataset

We have performed two tests using data from the same dataset to evaluate the accuracy of each model. In the first test, we used 60% of the original dataset, which was not used for training. In the second test, we used 100% of the original dataset, which contains 40% of the


data that is used for training the model itself. We expect that the second test for any model will outperform the first test because it contains data that is used for training.

### GaussianNB

Accuracy with test data: **0.8007414272474513**

Accuracy with the numbers [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]: **0.8408458542014469**

The Gaussian Naive Bayes model is the simplest model among the three. Its relatively low accuracy suggests that the pixel intensities in the digit data set do not follow a Gaussian distribution very strongly. Yet, it is still a decent baseline with over 80% accuracy. The results for the two tests were as expected, since the second test included some training data and had higher accuracy than the first test.




### KNN

Accuracy with test data: **0.9545875810936052**

Accuracy with the numbers [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]: **0.9677239844184753**

The K-Nearest Neighbors performs the best among the three models. The model's high accuracy suggests that similar digits in pixel space are clustered closely, and it effectively captures the local structure of the data. The results for the two tests were as expected, since the second test included some training data and had higher accuracy than the first test.




## MLP

Accuracy with test data: **0.9147358665430955**

Accuracy with the numbers [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]: **0.9488035614913745**

The Multilayer Perceptron performs better than GaussianNB but not as well as KNN. Although the MLP model is more flexible and powerful, it did not outperform KNN in this experiment because the number of training iterations was relatively low. The MLP did not have enough time to fully converge and demonstrate its full potential. The results for the two tests were as expected, since the second test included some training data and had higher accuracy than the first test.



## Poison Dataset

The poison in this project is represented by: on lines 129, 130

```
129  noise_scale = 10.0
130  poison = rng.normal(scale=noise_scale, size=X_train.shape)
```


We set poison to be a 3D nparray, size of samples, from a normal distribution of mean=0 and std=noise\_scale.

Then: on line 132

```
132  X_train_poison = X_train + poison
```

We add them up to mess with original intensity pixel data.

We get the result looks like picture below:



## Denoised Dataset

Kernel PCA is used in this project: on lines 202-209

```
202  kpca = KernelPCA(
203      n_components=None,
204      kernel='rbf',
205      gamma=0.001,
206      fit_inverse_transform=True,
207      alpha=5e-3,
208      random_state=42
209  )
```

Kernel PCA (Principal Component Analysis) is able to explore high-dimensional features, instead of only linear, and the RBF (Radial Basis Function) kernel is used here.

$\gamma$  parameter:

$$k(x, y) = e^{(-\gamma \|x - y\|^2)}$$

According to the formula, we can know that  $\gamma$  is a parameter to control the distance of the inputs, and small  $\gamma$  means underfitting, linear-like behavior, whereas big  $\gamma$  means overfitting.

$\alpha$  parameter:


This is a ridge regression, L2 penalty for RBF, in this case.

According to the lecture for regularization, we know that:

$$\min_{\beta} ||X\beta - y||_2^2 + \lambda ||\beta||_2^2$$

In this case, we don't use linear regression, but with matrix for RBF and KPCA feature coefficients for inverse\_transform, and we use  $\alpha$  to represent  $\lambda$  for ridge penalty here. The effect of it is just like to smooth out the image, and makes it easier to identify.

We get the result looks like picture below:




## Results on Poisoned/Denoised Dataset

### Poisoned Dataset

Table 1:

Model/Data	Original on Test	Poison on Test	Original on Num	Poison on Num	Denoise on Test	Denoise on Num
Gaussian	0.8007414272 474513	0.81464318 81371641	0.840845854 2014469	0.8319421 257651641	0.8489341 983317887	0.86533110 74012242
KNN	0.9545875810 936052	0.60518999 07321594	0.967723984 4184753	0.6505286 588759043	0.8266913 809082483	0.8619922 092376182
MLP	0.9147358665 430955	0.79425393 8832252	0.948803561 4913745	0.8046744 574290484	0.8257645 968489342	0.8553144 129104062

Poisoned Gaussian: [:200]



Poisoned KNN: [:200]

0 1 1 8 4 8 6 7 8 0 0 1 2 3 4 5 6 0 8 3 0 1 2 3 4 5 6 5 5 8 0 9 5 5 6 5 0 0 8 9 8 4 1 7 7 3 4 1 0 0 7 0 7 8 4 0 1 1 6 3 3 8 9 3 4 6 6 6 4 8 4 5 0 3 5 2 8 1 0 0

1 7 6 3 1 8 7 4 6 3 3 3 0 1 7 8 8 4 3 4 4 0 5 3 6 9 1 1 7 5 4 4 7 2 8 0 0 5 7 9 5 4 8 8 4 0 0 8 0 8 0 9 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 0 0 1 2 3 4 5 6 7 8 8

0 3 5 5 6 5 0 9 8 8 4 4 1 7 7 3 5 9 0 0 2 2 7 8 6 0 1 8 6 3 3 7 0 3 4 6 6 6 4 8

Poisoned MLP: [:200]

0 1 8 9 4 9 6 7 8 9 0 1 2 3 4 5 6 8 8 9 0 1 2 3 1 5 6 1 8 9 0 9 5 5 6 5 0 9 8 9 8 4 1 7 7 3 5 1 0 0 8 1 7 8 1 0 1 1 6 3 3 8 3 3 4 6 6 6 4 8 1 5 0 9 8 1 8 1 0 0

1 7 6 3 2 1 7 4 6 3 1 3 9 1 7 1 8 4 3 1 4 0 5 7 6 9 1 1 7 5 4 4 7 2 8 8 8 5 7 9 9 4 8 8 4 9 0 8 9 8 0 9 2 3 1 5 6 7 8 9 0 1 8 3 1 5 6 7 8 9 0 1 2 3 1 5 6 7 8 9

0 9 5 5 6 5 0 9 8 9 1 1 1 7 7 3 5 9 0 0 2 2 7 8 2 0 1 2 6 3 3 1 3 3 4 6 6 6 1 9

Denoised Gaussian: [:200]

0 1 1 3 4 9 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 9 5 5 6 5 0 3 5 9 8 4 1 7 7 3 5 1 0 0 2 1 7 7 7 0 1 1 6 3 3 7 3 3 4 6 6 6 4 7 1 5 0 9 5 1 2 1 0 0

1 7 6 3 2 1 7 4 6 3 1 3 9 1 7 6 8 4 3 1 4 0 5 7 6 9 6 1 7 5 4 4 7 2 8 8 2 5 7 9 5 4 8 8 4 3 0 8 9 8 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9

0 9 5 5 6 5 0 9 8 9 2 4 1 7 7 3 5 1 0 0 2 2 7 8 2 0 1 2 6 3 3 4 3 3 4 6 6 6 1 9

Denoised KNN: [:200]

0 1 1 3 4 9 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 1 5 6 7 8 9 0 9 5 5 6 5 0 3 7 9 8 4 1 7 7 3 5 1 0 0 2 2 7 7 2 0 1 1 6 3 3 7 3 3 4 6 6 6 4 7 1 5 0 9 5 2 8 1 0 0

1 7 6 3 2 1 7 4 6 3 1 3 9 1 7 1 1 4 3 1 4 0 5 3 6 9 1 1 7 5 4 4 7 2 8 7 2 5 7 9 5 4 8 8 4 3 0 8 9 2 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9

0 9 5 5 6 5 0 9 8 9 8 4 1 7 7 3 5 1 0 0 2 2 7 8 1 0 1 2 6 3 3 7 3 3 4 6 6 6 4 9

Denoised MLP: [:200]

The overall results of the Gaussian model poison is 0.8146431881371641



The overall results of the Gaussian model poison with the numbers [0, 1, 2, 3, 4, 5, 6, 7, 8, 9] is 0.8319421257651641

Model 2 poison trained...

The overall results of the KNN model poison is 0.6051899907321594

The overall results of the KNN model poison with the numbers [0, 1, 2, 3, 4, 5, 6, 7, 8, 9] is 0.6505286588759043

Model 3 poison trained...

The overall results of the MLP model poison is 0.794253938832252

The overall results of the MLP model poison with the numbers [0, 1, 2, 3, 4, 5, 6, 7, 8, 9] is 0.8046744574290484

Samples of poisoned training sets:

Some denoised training samples:

Model 1 denoised trained...

The overall results of the Gaussian model denoised is 0.8489341983317887

The overall results of the Gaussian model denoised with the numbers [0, 1, 2, 3, 4, 5, 6, 7, 8, 9] is 0.8653311074012242

Model 2 denoised trained...

The overall results of the KNN model denoised is 0.8266913809082483

The overall results of the KNN model denoised with the numbers [0, 1, 2, 3, 4, 5, 6, 7, 8, 9] is 0.8619922092376182

Model 3 denoised trained...

The overall results of the MLP model denoised is 0.8257645968489342

The overall results of the MLP model denoised with the numbers [0, 1, 2, 3, 4, 5, 6, 7, 8, 9] is 0.8553144129104062